


COMPARATIVE ANALYSIS OF CONVOLUTION OPERATIONS, REGULARIZATION TECHNIQUES, AND CNN ARCHITECTURES

A Scenario-Based Study for Plant Disease Image Classification

Leaf Image → Convolutional Filters → Feature Maps



Field	Details
Student Name	Paleru Dharani Govardhan
Register No.	192525280
Department	Artificial Intelligence and Machine Learning
Course	Deep Learning
Bloom's Level	BL4 – Analyse
Assessment Tool	AT3 – Comparative Analysis

ABSTRACT

This report presents a comparative, analytical study (Bloom's Level BL4 – Analyse) of the core building blocks of **Convolutional Neural Networks (CNNs)** — convolution operations, regularization techniques, and landmark CNN architectures — in the context of **plant disease image classification**. Convolutional operations, including standard, strided, dilated, depthwise, pointwise, and depthwise-separable convolutions, are examined for their working principles, computational cost, and suitability for leaf-image feature extraction. Regularization techniques — dropout, L1/L2 penalties, data augmentation, batch normalization, and early stopping — are analysed for their role in controlling overfitting, particularly under the small-dataset conditions typical of agricultural image datasets. Four influential CNN architectures, AlexNet, VGG-16, ResNet-50, and DenseNet-121, are compared along dimensions of depth, parameter count, computational cost, and gradient-flow behaviour. Building on this analysis, a scenario-based recommendation is developed for a plant-disease classification system that must operate on a relatively small, visually inconsistent dataset: a combination of depthwise-separable convolution, data augmentation with dropout and early stopping, and a DenseNet/ResNet-style backbone is proposed and justified. The report closes with a critical-analysis section addressing why deeper networks do not always improve performance, how filter size trades off receptive field against computation, and which architecture is best suited to this scenario, supported throughout by mathematical formulations, worked numerical examples, comparative tables, and technical diagrams.

1. INTRODUCTION

Convolutional Neural Networks are a class of deep learning models specifically designed to process data with a grid-like topology, most commonly images. Unlike fully connected networks, CNNs exploit **local connectivity** and **weight sharing**: a small kernel slides across the image and applies the same learned weights at every spatial location. This design drastically reduces the number of parameters compared to a fully connected layer operating on raw pixels, while preserving the spatial structure of the image — an essential property when the object of interest (for example, a diseased region on a leaf) can appear anywhere in the frame.

CNNs are powerful for image classification because they build a **hierarchical representation** of visual information. Early convolutional layers respond to simple, low-level patterns such as edges and colour gradients; intermediate layers combine these into textures and parts; and deep layers assemble parts into whole-object or whole-pattern representations. The convolution operation itself is what allows this hierarchy to emerge: by computing a weighted sum over a local neighbourhood of pixels, a convolutional filter extracts a specific type of spatial feature (an edge orientation, a texture pattern, a colour transition) and produces a feature map indicating where in the image that feature occurs.

Different **convolution variants** exist because a single 3×3 sliding-window operation is not equally efficient or equally effective for every task. Strided convolutions trade spatial resolution for reduced computation; dilated convolutions expand the receptive field without adding parameters; depthwise and pointwise convolutions factorize the standard convolution to sharply cut computational cost, which matters for deployment on resource-constrained or mobile devices. Selecting the right convolution operation is therefore an engineering trade-off between accuracy, receptive field, and computational budget.

Regularization is required because CNNs are heavily over-parameterized relative to the amount of labelled training data usually available, especially in specialized domains such as agriculture. Without regularization, a network can memorise the training set — including irrelevant background detail, lighting artefacts, and noise — rather than learning generalisable disease-indicative features. Techniques such as dropout, weight penalties, data augmentation, batch normalization, and early stopping each attack this problem from a different angle, and, as this report will show, they are most effective when combined rather than used in isolation.

Finally, **CNN architecture selection** matters because the way convolutional and pooling layers are organised (depth, connectivity pattern, filter sizes) directly determines how well gradients propagate during training, how efficiently parameters are used, and how well the network generalises from limited data. AlexNet, VGGNet, ResNet, and DenseNet each represent a distinct architectural philosophy, and the choice among them should be driven by the specific constraints of the deployment scenario rather than by raw benchmark accuracy alone.




Figure 1: General CNN Image Classification Pipeline — Input Leaf Image → Convolution → Feature Maps → Pooling → Deep Features → Classifier → Disease Class.

2. CNN FUNDAMENTALS

Before comparing individual techniques, it is necessary to establish the fundamental CNN processing pipeline that every architecture in this report builds upon.

- 1. Input Image** — A raw leaf image, typically resized and normalised to a fixed resolution (e.g., 224×224×3) before entering the network.
- 2. Convolution** — Learnable kernels slide across the image, computing local weighted sums that highlight specific spatial patterns.
- 3. Activation** — A non-linear function (commonly ReLU) is applied element-wise to the convolution output, allowing the network to model non-linear relationships.
- 4. Feature Map** — The activated output of a convolutional layer — a 2-D grid indicating the strength of a detected feature at every spatial location.
- 5. Pooling** — A down-sampling operation (max or average pooling) that reduces spatial dimensions while retaining the most salient activations.
- 6. Deeper Convolution Layers** — Successive convolution + activation + pooling stages progressively build more abstract, higher-level feature representations.
- 7. Flattening / Global Pooling** — The final spatial feature maps are converted into a single feature vector, either by flattening or by global average pooling.
- 8. Fully Connected Layer** — One or more dense layers combine the extracted global features to model relationships relevant to the classification task.
- 9. Classification** — A final softmax layer converts the network's output into class probabilities — e.g., Healthy, Leaf Spot, Rust, Blight, or Powdery Mildew.

Hierarchical Feature Learning

A central analytical insight in CNN design is that **depth creates a feature hierarchy**. **Early layers** operate directly on pixel neighbourhoods and therefore detect low-level, generic patterns such as edges, colour boundaries, and simple textures. **Middle layers** combine these primitives into more complex, still fairly generic patterns — corners, curves, repeated textures. **Deep layers** combine middle-layer features into complex, task-specific structures that are far more discriminative for the classification target.

For plant disease detection specifically, this hierarchy can be interpreted as: **Edges → Veins → Spots → Lesions → Disease-specific patterns**. Early layers detect the edges bounding the leaf and its veins; middle layers begin to isolate localized spots or discolouration; deep layers combine spot shape, colour, distribution, and boundary characteristics into representations that are specific to a particular disease category (e.g., the concentric-ring pattern typical of certain leaf-spot diseases versus the powdery surface texture of powdery mildew).



Figure 2: Feature Hierarchy Learned by CNN Layers for Leaf Images — from generic edges to disease-specific patterns.

The Artificial Neuron

At the computational core of every convolutional filter and every fully connected layer lies the artificial neuron, which computes a weighted sum of its inputs, adds a bias term, and passes the result through a non-linear activation function:

$$z = \sum_i w_i x_i + b, \quad \text{output} = f(z)$$




Figure 3: Structure of an Artificial Neuron — weighted inputs, summation, bias, and non-linear activation.

PART A — CONVOLUTION OPERATIONS AND FILTERS

This part analyses six convolution operations relevant to CNN-based leaf-image classification: standard, strided, dilated (atrous), depthwise, pointwise, and depthwise-separable convolution. Each is examined for its working principle, mathematical formulation, computational behaviour, and suitability for the plant-disease scenario.

A.1 Standard Convolution

Standard convolution slides a learnable $K \times K$ kernel across the input, computing an element-wise multiplication between the kernel weights and the underlying pixel neighbourhood, followed by a summation. This weighted sum becomes one entry of the output feature map. The same kernel weights are reused at every spatial position (weight sharing), which allows the operation to detect a given pattern regardless of where it appears in the image.




Figure 4: Standard Convolution Operation — a 3×3 kernel slides over a 6×6 input to produce a 4×4 feature map.

A small numerical example clarifies the computation. Consider a 5×5 input and a 3×3 edge-detecting kernel; the value at the top-left output position is obtained by multiplying the overlapping 3×3 input patch element-wise with the kernel and summing the results:




Figure 5: Worked Numerical Example of a 3×3 Convolution on a 5×5 Input.

The spatial size of the output feature map is governed by the standard convolution output-size formula:

$$\text{Output} = \text{floor}((N + 2P - K) / S) + 1$$

where N is the input size, P is the padding, K is the kernel size, and S is the stride. For N = 5, P = 0, K = 3, S = 1, the output size is $\text{floor}((5 - 3)/1) + 1 = 3$, matching the 3×3 feature map above.

Computational Complexity

For an input of size H×W with C_{in} channels, a standard convolutional layer producing C_{out} output channels using a K×K kernel requires approximately $K^2 \times C_{in} \times C_{out} \times H \times W$ multiply-accumulate operations. This cost grows multiplicatively with kernel size, channel depth, and spatial resolution. For high-resolution leaf images (e.g., 512×512) processed through layers with hundreds of channels, this cost becomes substantial, motivating the more efficient factorized operations analysed later in this section.

A.2 Strided Convolution

Stride defines the step size by which the kernel moves across the input. With **stride = 1**, the kernel visits every possible position, producing a densely sampled feature map. With **stride = 2** (or higher), the kernel skips positions, visiting roughly a quarter as many locations for stride 2 in two dimensions, which both down-samples the spatial dimensions and reduces computation in subsequent layers.




Figure 6: Effect of Stride on Kernel Sampling — Stride 1 (dense, 4×4 output) vs Stride 2 (sparse, 2×2 output) on a 6×6 input with a 3×3 kernel.

Using the output-size formula with $N = 6$, $P = 0$, $K = 3$: at stride 1 the output is $\text{floor}((6-3)/1)+1 = 4$ (a 4×4 feature map); at stride 2 the output is $\text{floor}((6-3)/2)+1 = 2$ (a 2×2 feature map).

Advantages

- Reduces spatial dimensions directly within the convolution, avoiding a separate pooling step.
- Reduces the computational burden on all subsequent layers, since they operate on a smaller feature map.
- Can replace pooling layers entirely in modern lightweight architectures.

Limitations

- Aggressive striding can discard fine spatial detail before the network has extracted it.
- Small or thin disease features (e.g., early-stage leaf spots) may be skipped over or blurred across neighbouring positions.

In the plant-disease scenario, moderate striding (e.g., stride 2 in early layers) is generally safe because disease symptoms typically occupy multiple pixels; however, excessive striding in the very first layers risks losing subtle early-stage lesions before the network can learn to represent them.

A.3 Dilated (Atrous) Convolution

Dilated convolution introduces gaps between kernel elements, controlled by a dilation rate d . A dilation rate of 1 corresponds to a standard convolution; a dilation rate of 2 spaces the kernel taps two pixels apart, expanding the area of the input the kernel sees without increasing the number of learnable parameters.




Figure 7: Normal 3×3 Convolution (dilation = 1, receptive field 3×3) vs Dilated 3×3 Convolution (dilation = 2, receptive field 5×5).

The effective kernel size (the span of input pixels the dilated kernel covers) is given by:

$$K_{eff} = K + (K - 1)(d - 1)$$

For $K = 3$ and $d = 2$: $K_{eff} = 3 + (2)(1) = 5$, so a 3×3 kernel with dilation 2 covers the same spatial extent as a standard 5×5 kernel, but with only 9 weights instead of 25.

Dilated convolution is useful in the plant-disease scenario because disease patterns manifest at different spatial scales: a small isolated spot requires fine local detail, while a spreading blight pattern requires context from a wider area of the leaf. Stacking convolutions with different dilation rates lets the network capture both scales without a proportional increase in parameters or a loss of resolution from aggressive pooling.

Advantages

- Expands the receptive field rapidly without increasing parameter count.
- Preserves spatial resolution, since no down-sampling is required to enlarge context.
- Enables multi-scale pattern capture when different dilation rates are combined.

Limitations

- Can introduce a 'gridding' artefact where certain input pixels are never sampled together.
- Very large dilation rates may cause the kernel to skip over small, localized features entirely.

A.4 Filter Size and Receptive Field

Filter (kernel) size directly determines the receptive field of a single convolutional operation, and therefore both the granularity of the feature it can detect and its computational cost.




Figure 8: Receptive Field Comparison — 1×1 , 3×3 , and 5×5 kernels on the same input grid.

Property	1×1	3×3	5×5
Receptive field	Single pixel	Local 3×3 neighbourhood	Wider 5×5 neighbourhood
Parameters (per Cin,Cout pair)	1	9	25
Typical use	Channel mixing (pointwise)	General-purpose feature extraction	Coarser, larger-scale patterns
Feature type	No spatial context	Fine-grained (edges, texture)	Coarser (broad regions, blotches)

A 1×1 kernel has no spatial receptive field and only recombines channel information at a single location (this is exploited in pointwise convolution, A.6). A 3×3 kernel is the most common choice in modern CNNs because two stacked 3×3 layers achieve the same 5×5 receptive field with fewer parameters ($2\times 9=18$ vs 25) and an extra non-linearity, which is why VGGNet and later architectures standardised on small filters.

A.5 Depthwise Convolution

Depthwise convolution applies a single filter per input channel independently — it does not mix information across channels. Given C_{in} input channels, a depthwise convolution uses C_{in} separate $K\times K$ filters (one per channel), producing C_{in} output feature maps.




Figure 9: Depthwise Convolution — each colour channel is filtered independently by its own kernel; there is no cross-channel mixing.

A.6 Pointwise Convolution

Pointwise convolution uses a 1×1 kernel applied across all channels at each spatial location. It performs no spatial filtering but linearly combines the information present across channels at that pixel, effectively acting as a per-pixel fully connected layer over the channel dimension.





Figure 10: Pointwise (1×1) Convolution — channel information is linearly combined at each spatial location.

A.7 Depthwise Separable Convolution

Depthwise separable convolution factorizes a standard convolution into two sequential, cheaper operations: a **depthwise convolution** (spatial filtering per channel) followed by a **pointwise convolution** (channel mixing). This factorization is the basis of lightweight CNN architectures designed for mobile and edge deployment.

Standard Convolution vs Depthwise Separable Convolution



Depthwise separable convolution factorizes spatial filtering (depthwise) and channel mixing (pointwise), sharply reducing multiply-add operations.

$$\text{Standard cost} \approx K^2 \cdot C_{in} \cdot C_{out} \cdot H \cdot W$$

$$\text{Depthwise-Separable cost} \approx K^2 \cdot C_{in} \cdot H \cdot W + C_{in} \cdot C_{out} \cdot H \cdot W$$

Figure 11: Standard Convolution vs Depthwise Separable Convolution — factorization into depthwise and pointwise stages.

The computational cost of each approach can be compared directly:

$$\text{Standard Convolution} \approx K \times K \times C_{in} \times C_{out} \times H \times W$$

$$\text{Depthwise Separable} \approx K \times K \times C_{in} \times H \times W + C_{in} \times C_{out} \times H \times W$$

Consider a numerical example with $K = 3$, $C_{in} = 64$, $C_{out} = 128$, and a 112×112 feature map (typical of an early-to-mid CNN stage). Standard convolution requires $3^2 \times 64 \times 128 \times 112 \times 112 \approx 924$ million multiply-adds. Depthwise separable convolution requires $(3^2 \times 64 \times 112 \times 112) + (64 \times 128 \times 112 \times 112) \approx 7.2$ million + 102.8 million ≈ 110 million multiply-adds — roughly an **8-fold reduction** in computation for this configuration.




Figure 12: Computational Cost Comparison — Standard Convolution vs Depthwise Separable Convolution for $K=3$, $C_{in}=64$, $C_{out}=128$, $H=W=112$.

This efficiency makes depthwise separable convolution particularly attractive for *lightweight CNNs* intended for mobile or edge deployment — for instance, a field-deployable smartphone application for on-the-spot plant disease diagnosis, where computational resources and battery life are limited.

A.8 Comparative Analysis of Convolution Operations

Operation	Kernel	Spatial Effect	Feature Extraction	Computational Cost	Main Advantage	Limitation	Suitable Application
Standard Convolution	$K \times K$, full Cin	Preserves resolution (with padding)	Strong, general-purpose	High ($K^2 \cdot \text{Cin} \cdot \text{Cout} \cdot H \cdot W$)	Rich feature representation	Computationally expensive at scale	General feature extraction backbone
Strided Convolution	$K \times K$, stride ≥ 2	Down-samples output	Coarser, fewer positions	Reduced (fewer output positions)	Learned down-sampling, no separate pooling	May skip small features	Replacing pooling; early down-sampling
Dilated Convolution	$K \times K$, gaps of $d-1$	Preserves resolution	Wider context, same detail	Same as standard for given K	Expands receptive field free of cost	Gridding artefacts at high d	Multi-scale context (e.g., blight spread)
Depthwise Convolution	$K \times K$ per channel	Preserves resolution	Per-channel spatial patterns only	Very low ($K^2 \cdot \text{Cin} \cdot H \cdot W$)	Sharp reduction in computation	No cross-channel mixing	First stage of lightweight blocks
Pointwise Convolution	1×1 , all channels	No spatial change	Cross-channel combination only	Low ($\text{Cin} \cdot \text{Cout} \cdot H \cdot W$)	Cheap channel mixing	No spatial receptive field	Second stage of lightweight blocks
Depthwise Separable Conv.	Depthwise + Pointwise	Preserves resolution	Near-standard quality, factorized	Low (sum of both stages)	Large computation reduction, few parameters	Slightly lower capacity than standard conv.	Mobile / edge CNNs (e.g., MobileNet)
Max Pooling	$K \times K$, no weights	Down-samples output	Retains strongest activation	Negligible	Translation robustness, cheap down-sampling	Discards non-maximal information	Down-sampling after conv. blocks
Average Pooling	$K \times K$, no weights	Down-samples output	Smooths activations	Negligible	Preserves overall activation level	Dilutes sharp, localized features	Global summarization (e.g., GAP before classifier)

Critical analysis. No single convolution operation is universally optimal. Standard convolution offers the strongest raw representational power and remains the default choice where computational budget is not the primary constraint. Strided convolution is preferable when spatial down-sampling is required as part of feature extraction (rather than as a separate pooling step). Dilated convolution is the right choice when a wider receptive field is needed without losing resolution — relevant when disease symptoms of very different spatial scale must be captured within the same network. Depthwise, pointwise, and their combination (depthwise separable convolution) are preferable whenever computational or memory budget is constrained, such as on mobile or embedded platforms for in-field disease diagnosis, at a modest cost in representational capacity that is often recovered by using slightly more channels or a deeper network at the same computational budget.

PART B — REGULARIZATION TECHNIQUES

This part compares five regularization techniques — dropout, L1/L2 regularization, data augmentation, batch normalization, and early stopping — and analyses the relationship between underfitting, good generalization, and overfitting.

B.1 Dropout

Dropout regularizes a network by randomly deactivating a fraction of neurons during each training step, forcing the remaining neurons to learn features that are useful independently of any particular subset of co-activated units. During inference, all neurons are active but their outputs are scaled to match the expected activation level seen during training. This **reduces co-adaptation** between neurons — the tendency of units to rely on the presence of specific other units — which is a major driver of overfitting.




Figure 13: Normal Network (all neurons active) vs Network with Dropout (a random subset of neurons deactivated per training step).

B.2 L1 and L2 Regularization

Both L1 and L2 regularization discourage large weight magnitudes by adding a penalty term to the loss function, but they induce different behaviours in the learned weights.

L1 Regularization

$$Loss = Original\ Loss + \lambda \sum |w|$$

The L1 penalty grows linearly with weight magnitude and has a constant gradient regardless of how small a weight becomes, which tends to push many weights exactly to zero. This produces sparse models where irrelevant input features effectively receive zero weight.

L2 Regularization

$$Loss = Original\ Loss + \lambda \sum w^2$$

The L2 penalty grows quadratically, so its gradient shrinks as weights approach zero; this produces weight decay — a smooth shrinkage of all weights toward small values rather than exact zeros — yielding smoother, more evenly distributed weight solutions that are less sensitive to any single input feature.




Figure 14: L1 Regularization (diamond-shaped constraint region, promotes sparsity) vs L2 Regularization (circular constraint region, promotes smooth shrinkage). The orange point marks the constrained solution; the red cross marks the unconstrained loss minimum.

B.3 Data Augmentation

Data augmentation artificially expands the effective size and diversity of the training set by applying label-preserving transformations to existing images, without collecting new data. For the plant-disease scenario, this is especially valuable because labelled disease images are relatively scarce and expensive to collect.




Figure 15: Data Augmentation Chain — Original Leaf Image → Rotation → Flip → Zoom → Brightness Change → Crop → Augmented Dataset.

Rotation and flipping address orientation variation, since a diseased leaf may be photographed at any angle. Brightness and contrast adjustments simulate lighting variation across different capture conditions (outdoor sunlight vs indoor/shaded imaging). Zoom and cropping simulate variation in camera distance and framing. Because these transformations preserve the disease label while varying nuisance factors, the network is encouraged to learn disease-indicative features rather than incidental capture conditions.

B.4 Batch Normalization

Batch normalization normalizes the activations of a layer across each mini-batch to have zero mean and unit variance, then applies a learned scale (γ) and shift (β) to restore representational flexibility:

$$\hat{x} = (x - \mu_B) / \sqrt{\sigma_B^2 + \epsilon}, \quad y = \gamma \hat{x} + \beta$$

Stabilizes distributions across mini-batches → faster, smoother convergence



$$\hat{x} = \frac{x - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}, \quad y = \gamma \hat{x} + \beta$$

Figure 16: Batch Normalization — activations are normalized per mini-batch, then rescaled and shifted by learned parameters.

By keeping the distribution of layer inputs stable throughout training, batch normalization mitigates internal covariate shift, allows the use of higher learning rates, improves gradient flow through deep networks, and — as a secondary effect — provides a mild regularizing influence due to the noise introduced by per-batch statistics.

B.5 Early Stopping

Early stopping monitors validation loss (or another validation metric) during training and halts training once the validation metric stops improving, even if the training metric continues to improve. This directly prevents the later stages of training in which the network begins memorising training-set-specific noise rather than learning generalisable patterns.




Figure 17: Early Stopping — Training loss keeps falling while validation loss begins rising after approximately epoch 18, marking the ideal stopping point.

B.6 Underfitting, Good Generalization, and Overfitting

Scenario	Training Performance	Validation Performance	Typical Cause
Underfitting	Poor	Poor	Model capacity too low, or insufficient training
Good Generalization	High	High and similar to training	Adequate capacity, appropriate regularization
Overfitting	Very high	Poor / deteriorating over epochs	Model memorises training data; regularization insufficient




Figure 18: Training vs Validation Accuracy and Loss across Underfitting, Good Generalization, and Overfitting scenarios.

On a small plant-disease dataset, the overfitting curve is the most likely default outcome if no regularization is applied: with only a few hundred or thousand labelled images and a deep CNN containing millions of parameters, the network has more than enough capacity to memorise the training set, causing training accuracy to approach 100% while validation accuracy plateaus or declines. This is precisely the condition that motivates the combined use of data augmentation, dropout, and early stopping recommended later in Part D.

PART C — CNN ARCHITECTURES

This part compares four landmark CNN architectures — AlexNet, VGG-16, ResNet-50, and DenseNet-121 — analysing how their architectural decisions affect accuracy potential, parameter count, computation, gradient flow, feature reuse, training difficulty, and generalization.

C.1 AlexNet

AlexNet (2012) was the architecture that demonstrated deep CNNs could dramatically outperform traditional computer-vision pipelines on large-scale image classification. It consists of five convolutional layers (the first using large 11×11 filters with stride 4, later layers using smaller filters) followed by three fully connected layers. Its historical importance lies in popularising **ReLU activation** (faster training than sigmoid/tanh), **dropout** in the fully connected layers to control overfitting, and GPU-based training, which made deep CNNs computationally practical for the first time.

C.2 VGGNet

VGGNet (2014) simplified architectural design by using only small **3×3 convolutional filters** stacked repeatedly, interspersed with 2×2 max-pooling layers, and increasing channel depth after each pooling stage. Stacking multiple 3×3 layers achieves a receptive field equivalent to a single larger filter (e.g., two 3×3 layers $\approx$ one 5×5 layer) while using fewer parameters and adding extra non-linearities. This simple, uniform design made VGGNet easy to understand and extend, but its depth and fully connected layers give it a very large parameter count.




Figure 19: AlexNet Architecture — 5 convolutional layers followed by 3 fully connected layers.

VGG-16 — 13 Convolutional Layers (all 3×3) + 3 Fully Connected Layers




Figure 20: VGG-16 Architecture — 13 convolutional layers (all 3×3) organised into 5 blocks, followed by 3 fully connected layers.

Aspect	AlexNet	VGG-16
Depth	8 weight layers (5 conv + 3 FC)	16 weight layers (13 conv + 3 FC)
Filter sizes	11×11, 5×5, 3×3 (mixed)	3×3 throughout
Parameters (approx.)	~60 million	~138 million
Computation	Lower than VGG (fewer, larger-stride layers)	High (many 3×3 conv. layers at full resolution)
Feature extraction	Coarser, less refined hierarchical features	Finer hierarchical features via repeated small filters
Advantages	Proof of concept for deep CNNs; relatively fast	Simple, uniform design; strong transferable features
Limitations	Large filters lose fine detail early; less depth	Very high parameter count and memory cost

C.3 ResNet

As networks are made deeper, training accuracy can paradoxically *degrade* rather than improve — not due to overfitting, but because very deep plain networks become difficult to optimise (the **degradation problem**), compounded by the **vanishing gradient problem**, where gradients shrink as they are back-propagated through many layers, leaving early layers barely updated.

ResNet (2015) addresses this with **residual blocks**: instead of learning a direct mapping $H(x)$, each block learns a residual function $F(x)$ and adds it back to the input via a **skip (identity) connection**:




Figure 21: Traditional (Plain) Block vs Residual Block in ResNet — the skip connection adds the identity of x back to the transformed output $F(x)$.

This formulation makes it easy for a block to learn an identity mapping (simply drive $F(x)$ toward zero) when additional depth is not helpful, which means adding more layers can never make optimisation harder than the shallower network

— it can only help. The skip connection also provides a direct path for gradients to flow backward without passing through every intermediate non-linearity, substantially alleviating vanishing gradients and enabling stable training of networks with 50, 101, or more layers.

C.4 DenseNet

DenseNet (2017) extends the skip-connection idea further with **dense connectivity**: every layer within a dense block receives, as input, the *concatenation* of the feature maps of all preceding layers in that block, rather than only the output of the immediately preceding layer:

$$x_l = H_l([x_0, x_1, \dots, x_{l-1}]) \quad \text{where } [] \text{ denotes concatenation}$$




Figure 22: Dense Connectivity in DenseNet — every layer receives feature maps from all preceding layers via concatenation (three representative skip paths shown).

Because early features remain directly accessible to every later layer, DenseNet strongly encourages **feature reuse**: later layers do not need to re-learn features already computed earlier, which allows DenseNet to achieve competitive accuracy with substantially fewer parameters than a comparably deep ResNet or VGGNet. The dense connectivity pattern also provides deep supervision, since gradients from the loss can reach any layer directly through short paths, further improving gradient flow in very deep configurations.

C.5 Complete Architecture Comparison

The following comparison uses one representative, historically accurate variant of each architecture — AlexNet, VGG-16, ResNet-50, and DenseNet-121 — so that reported figures are not mixed across variants of the same architecture family.

Parameter	AlexNet	VGG-16	ResNet-50	DenseNet-121
Year introduced	2012	2014	2015	2017
Basic architecture	Sequential conv + FC	Sequential conv blocks + FC	Residual blocks with skip connections	Dense blocks with concatenated connections

Parameter	AlexNet	VGG-16	ResNet-50	DenseNet-121
Number of layers	8 weight layers	16 weight layers	50 weight layers	121 weight layers
Convolutional filters	11×11, 5×5, 3×3	3×3 throughout	1×1 and 3×3 (bottleneck blocks)	1×1 and 3×3 (bottleneck, dense blocks)
Key innovation	ReLU + Dropout + GPU training	Small, uniform 3×3 filters, depth	Residual (skip) connections	Dense connectivity, feature reuse
Approx. parameter size	~60 million	~138 million	~25.6 million	~8 million
Computational requirement	Moderate	High	Moderate–High	Moderate (efficient per-parameter)
Vanishing/degradation handling	Not addressed	Not addressed	Solved via identity skip connections	Solved via dense skip connections
Feature reuse	Minimal	Minimal	Partial, via residual addition	Extensive, via concatenation
Advantages	Simple, historically foundational	Uniform design, strong transfer features	Enables very deep, stable training	Parameter-efficient, strong gradient flow
Limitations	Shallow by modern standards	Very high parameter/memory cost	Deeper variants still parameter-heavy	Higher memory use from concatenated features
Typical applications	Baseline / educational benchmark	Transfer learning, feature extraction	General-purpose deep vision backbones	Data-efficient, fine-grained classification



Figure 23: CNN Architecture Evolution Timeline — AlexNet → VGG-16 → ResNet-50 → DenseNet-121.

The evolution from AlexNet to DenseNet reflects a consistent design trend: as networks grew deeper to capture richer hierarchical features, connectivity patterns (skip connections, dense connections) became necessary to keep training stable and parameter-efficient. AlexNet proved that depth and GPU-scale training were viable; VGGNet showed that small, uniform filters stacked deeply could extract strong features, at the cost of parameter efficiency; ResNet solved the degradation problem that arises when depth increases further; and DenseNet pushed connectivity to its logical extreme, maximising feature reuse and parameter efficiency.

PART D — SCENARIO-BASED COMPARATIVE ANALYSIS

Scenario: Plant Disease Classification

A university wants to develop an image classification system to identify different types of plant diseases from leaf images. The available dataset is **relatively small**, and the images exhibit **lighting variation**, **orientation variation**, and **background variation**. This section applies the convolution, regularization, and architecture analysis from Parts A–C to derive and justify a complete, scenario-specific system design.

D.1 Problem Analysis

Small dataset — A CNN with millions of parameters trained on only a few hundred or thousand labelled leaf images is at high risk of overfitting: the network has more than enough capacity to memorise individual training images rather than learn generalisable disease-indicative patterns, leading to strong training performance but poor performance on new, unseen leaves.

Lighting variation — Images captured under different lighting conditions (direct sunlight, shade, indoor lighting) produce different pixel intensity distributions for the same underlying disease pattern. Without correction, the network may partly learn to associate lighting conditions with class labels rather than the disease itself.

Orientation variation — Leaves and disease patterns may appear at arbitrary rotation angles and orientations depending on how the photograph was taken. A network that has only ever seen upright examples of a disease pattern may fail to recognise the same pattern rotated by 90° or reflected.

Background variation — Leaf images may be captured against soil, other foliage, hands, or artificial backgrounds. If background diversity is not addressed, the network risks learning background cues (e.g., a particular soil colour correlating by chance with a particular disease class in the training set) instead of genuine disease features — a spurious-correlation risk that can severely hurt generalisation to new backgrounds.




Figure 24: Challenge → Solution Mapping for the Plant Disease Classification Scenario.

D.2 Convolution Operation Selection

The primary recommendation for the convolutional backbone is **depthwise separable convolution**, justified as follows:

- Lower computational cost — as shown in A.7, depthwise separable convolution can reduce multiply-add operations by a large factor relative to standard convolution for typical channel configurations, which matters if the system is intended to run in resource-constrained settings such as field devices or low-cost servers.
- Efficient feature extraction — despite the factorization, depthwise separable convolution retains most of the representational quality needed to distinguish visually distinct disease categories.
- Suitable for lightweight CNN systems — architectures built from depthwise separable blocks (e.g., MobileNet-style designs) are specifically intended for exactly this kind of deployment-constrained scenario.
- Reduced parameter count — fewer parameters lowers overfitting risk on a small dataset, complementing the regularization strategy in D.3.

Trade-off acknowledgement. Standard convolution may still provide stronger raw representational capacity in settings where computational budget is not constrained, and the choice between standard and depthwise separable convolution should ultimately be validated experimentally on held-out data rather than assumed from theory alone.

Complementary technique. Dilated convolution is worth incorporating as a complementary technique — not a replacement — in selected layers, because plant disease patterns vary in spatial scale (a small isolated spot versus a spreading blight region). A small dilation rate in mid-level layers can expand context without the resolution loss associated with additional pooling. This is distinguished as a *primary recommendation (depthwise separable convolution)* versus a *possible complementary technique (dilated convolution in selected layers)*.

D.3 Regularization Selection

A combination of techniques is recommended, since each addresses a different aspect of the small-dataset, high-variation problem:

1. Data Augmentation — Rotation, flipping, brightness/contrast adjustment, zoom, and cropping directly counteract lighting variation, orientation variation, and background/framing differences by exposing the network to many synthetic variations of each labelled example.

2. Dropout — Applied in the later, fully connected or classification-head layers, dropout reduces the co-adaptation of neurons and mitigates overfitting arising from the small dataset size.

3. Early Stopping — Training is halted once validation performance stops improving, preventing the network from continuing to fit training-set-specific noise once useful generalisable patterns have already been learned.

Optional — L2 Regularization: A weight-decay penalty can be added as an additional safeguard against large, over-confident weights, particularly useful if the network still shows signs of overfitting after augmentation, dropout, and early stopping are applied.

Combining multiple techniques is more appropriate than relying on a single one because each targets a distinct failure mode: augmentation increases effective data diversity *before* training, dropout constrains the network *during* training, and early stopping intervenes based on validation behaviour *throughout* training. No single technique addresses all three roles simultaneously, so their combination is complementary rather than redundant.

Small Dataset → Augmentation → CNN → Dropout / L2 → Early Stopping → Better Generalization

D.4 CNN Architecture Selection

For the plant-disease scenario, both **DenseNet** and **ResNet** are analysed before a final recommendation is made, since either could reasonably be justified depending on deployment priorities.

Factor	DenseNet	ResNet
Feature reuse	Extensive — every layer accesses all preceding feature maps directly, which is valuable when the training set is small and every learned feature should be maximally exploited.	Partial — features are combined via addition rather than concatenation, so reuse is less explicit.
Parameter efficiency	High — DenseNet-121 (~8M parameters) achieves competitive accuracy with far fewer parameters than ResNet-50 (~25.6M), reducing overfitting risk on small data.	Moderate — more parameters than DenseNet for comparable depth, increasing overfitting risk unless data or regularization compensates.
Gradient flow	Strong — dense skip connections give the loss short, direct paths to early layers.	Strong — identity skip connections also address vanishing gradients effectively.
Small-data behaviour	Favourable — parameter efficiency and feature reuse both directly reduce the risk of overfitting on limited labelled data.	Reasonable, but relatively more parameters make it slightly more prone to overfitting on very small datasets without strong regularization.
Fine-grained visual patterns	Well suited — reuse of low-level features (edges, colour transitions) benefits fine-grained distinctions between visually similar disease classes.	Also capable, though less explicit reuse of low-level detail in very deep stages.
Computational cost	Moderate — concatenation increases memory usage during training despite fewer parameters.	Moderate–High — more parameters but simpler additive connections, sometimes more memory-efficient at inference.
Transfer learning suitability	Strong — pretrained DenseNet backbones transfer well to small, specialized image domains such as agricultural imagery.	Strong — pretrained ResNet backbones are also widely used and well validated for transfer learning.

On balance, **DenseNet (specifically DenseNet-121)** is recommended as the primary backbone for this scenario, primarily because of its parameter efficiency and strong feature reuse, both of which directly reduce overfitting risk on a small, high-variation dataset. This is not a claim that DenseNet is universally superior — for a larger, more resource-abundant deployment with a substantially larger labelled dataset, ResNet-50 would remain a highly competitive and simpler-to-train alternative, and the final choice should be confirmed by empirical validation on the actual dataset.

D.5 Comparison with Alternatives

Factor	DenseNet (Selected)	ResNet	VGGNet
Accuracy potential	High, given sufficient training	High, given sufficient training	Moderate–High, but more parameters to fit
Parameter efficiency	High (~8M, DenseNet-121)	Moderate (~25.6M, ResNet-50)	Low (~138M, VGG-16)
Computational cost	Moderate (memory-heavy from concatenation)	Moderate–High	High
Feature reuse	Extensive (concatenation)	Partial (addition)	Minimal

Factor	DenseNet (Selected)	ResNet	VGGNet
Gradient flow	Strong (dense skip paths)	Strong (identity skip paths)	Weaker in very deep configurations
Small dataset suitability	Favourable	Reasonable	Higher overfitting risk
Transfer learning	Strong, widely used pretrained weights	Strong, widely used pretrained weights	Strong, but heavier fine-tuning cost
Overfitting risk	Lower, due to fewer parameters	Moderate	Higher, due to very high parameter count

Recommendation. DenseNet-121 is recommended as the primary backbone for the plant-disease classification system, with ResNet-50 identified as the strongest alternative should training-memory constraints (arising from DenseNet's feature concatenation) become limiting in practice.

D.6 Proposed System Architecture




Figure 25: Proposed Lightweight CNN System for Plant Disease Classification — from raw leaf image to disease class prediction. Disease classes shown (Healthy, Leaf Spot, Rust, Blight, Powdery Mildew) are illustrative examples.

D.7 Performance and Generalization Analysis

The proposed combination of techniques is expected to improve system behaviour along several dimensions, each traceable to a specific mechanism:

Technique	Mechanism	Expected Benefit
Data Augmentation	Exposes the network to many diverse synthetic variations of each labelled example	Better generalization to unseen lighting, orientation, and framing
Dropout	Reduces co-adaptation between neurons during training	Reduced overfitting on the small labelled dataset
Depthwise Separable Convolution	Factorizes spatial and channel-wise computation	Lower computational cost, feasible for constrained deployment

Technique	Mechanism	Expected Benefit
DenseNet Dense Connections	Provides direct gradient paths and feature reuse across layers	Easier optimisation and stronger performance with fewer parameters
Early Stopping	Halts training once validation performance stops improving	Avoids over-training on training-set-specific noise

Together, these mechanisms are expected to jointly improve classification accuracy (through stronger, reusable feature extraction), generalization (through augmentation and regularization), computational efficiency (through depthwise separable convolution), gradient flow and training stability (through dense connectivity), and robustness to real-world capture variation (through the combined effect of augmentation and a feature-reuse-oriented backbone).

D.8 Challenges and Solutions

Challenge	Solution
Small Dataset	Data augmentation, transfer learning from pretrained weights, and regularization (dropout, early stopping) collectively compensate for limited labelled examples.
Background Bias	Background diversification during data collection where possible, combined with cropping/segmentation and augmentation to reduce the network's reliance on background cues.
Lighting Variation	Brightness and contrast augmentation during training, combined with input normalization, to make the network robust to varying illumination.
Class Imbalance	Class weighting in the loss function, balanced sampling during training, and targeted augmentation of under-represented disease classes.

PART E — CRITICAL ANALYSIS

Q1. Why does increasing the number of convolutional layers not always improve CNN performance?

Adding convolutional layers increases model capacity, but capacity alone does not guarantee better performance, for several interacting reasons. First, **overfitting**: greater capacity allows the network to memorise training-set-specific detail rather than learn generalisable patterns, especially when the training dataset is small relative to parameter count. Second, **vanishing gradients**: in very deep plain (non-residual) networks, gradients shrink as they are back-propagated through many layers during training, leaving early layers under-trained. Third, the **degradation problem**: empirically, plain networks made deeper beyond a certain point show *higher* training error, not just validation error — indicating an optimisation difficulty rather than overfitting, since even training performance suffers. Fourth, **computational cost** grows with depth, increasing training time and resource requirements without a guaranteed proportional accuracy gain. Fifth, **optimisation difficulty** increases in deep, unstructured stacks of layers, as the loss landscape becomes harder to navigate. Finally, once a network is sufficiently deep for a given task and dataset, additional layers tend to produce **diminishing returns** — the marginal accuracy gain per additional layer decreases. ResNet's residual (skip) connections directly address the degradation problem: because a residual block can easily learn an identity mapping when extra depth is not needed, adding layers can no longer make optimisation strictly harder than the shallower baseline.

Q2. How do different filter sizes affect feature extraction and computational cost?

Filter size directly sets the receptive field of a single convolutional layer and its parameter/computation cost. A **1×1** filter has no spatial receptive field at all — it only recombines information across channels at a single pixel, at the lowest possible computational cost per output. A **3×3** filter captures a small local neighbourhood, extracting fine-grained features such as edges and textures, at moderate cost (9 weights per channel pair). A **5×5** filter captures a larger neighbourhood in a single operation, extracting coarser, broader patterns, at higher cost (25 weights per channel pair) — but two stacked 3×3 layers can achieve the same 5×5 receptive field with fewer total parameters (18 vs 25) and an additional non-linearity, which is why modern architectures favour small, stacked filters over single large filters.




Figure 26: Receptive Field Comparison — 1×1, 3×3, and 5×5 kernels illustrate how filter size determines spatial context per operation.

Q3. Why are depthwise separable convolutions computationally efficient?

Compare the parameter counts of standard convolution and depthwise separable convolution:

$$\text{Standard Convolution: } K^2 \times \text{Cin} \times \text{Cout}$$

$$\text{Depthwise Separable: } K^2 \times \text{Cin} + \text{Cin} \times \text{Cout}$$

Standard convolution scales *multiplicatively* with K^2 , Cin, and Cout together in a single term. Depthwise separable convolution splits this into an *additive* combination of two much smaller terms: a spatial term ($K^2 \times \text{Cin}$) that does not depend on Cout, and a channel-mixing term ($\text{Cin} \times \text{Cout}$) that does not depend on K^2 . As a numerical example, with $K = 3$, $\text{Cin} = 64$, $\text{Cout} = 128$: standard convolution requires $3^2 \times 64 \times 128 = 73,728$ weights, while depthwise separable convolution requires $(3^2 \times 64) + (64 \times 128) = 576 + 8,192 = 8,768$ weights — roughly an **8.4-fold reduction** in parameters (and a similar reduction in multiply-add operations per spatial position), because the expensive Cout-scaling only applies to the cheap 1×1 pointwise stage rather than to the full $K \times K$ spatial stage.

Q4. Why is regularization particularly important when training CNNs on small datasets?

CNNs typically contain a very **high parameter count** — often millions of learnable weights — relative to the number of labelled training images available in specialized domains such as plant pathology. When the dataset is small, the network has more than enough capacity to **memorise** individual training examples, including their incidental noise, rather than learning the generalisable visual patterns that define each class. This produces classic **overfitting**: training performance appears excellent while **validation performance** is poor, because the memorised details do not transfer to new images. **Data augmentation** artificially increases effective dataset diversity; **dropout** limits the network's ability to rely on brittle, co-adapted neuron combinations; **weight regularization** (L1/L2) discourages excessively large, over-confident weights; and **early stopping** halts training before the network transitions from learning generalisable patterns to memorising noise. Used together, these techniques compensate for the capacity–data mismatch inherent to training large CNNs on small, specialized datasets.

Q5. How does ResNet overcome the degradation problem?

ResNet reformulates what each block learns. Instead of directly learning a target mapping $H(x)$, a residual block learns the **residual function** $F(x) = H(x) - x$, and the block's actual output is computed by adding the input back via a skip connection:

$$y = F(x) + x$$

This reformulation is easier to optimise: if the optimal mapping for a given block is close to the identity, the network only needs to drive $F(x)$ toward zero — a much easier target than learning an exact identity mapping through a stack of non-linear layers. Skip connections also provide a **direct, largely unimpeded path for gradients** to flow from later layers back to earlier layers during backpropagation, without necessarily passing through every intervening non-linearity, which substantially reduces the vanishing-gradient effect and allows very deep networks (50+ layers) to train stably.




Figure 27: Residual Block — the identity skip connection allows gradients to bypass intermediate layers and allows the block to easily default to an identity mapping.

Q6. How does DenseNet promote feature reuse?

Within a dense block, every layer's input is the **concatenation** of the feature maps produced by *all* preceding layers in that block, not merely the output of the immediately previous layer:

$$x_l = H_l([x_0, x_1, \dots, x_{l-1}])$$

Because low-level features computed early in the block (edges, colour transitions) remain directly accessible, unmodified, to every subsequent layer, later layers do not need to re-derive these features from scratch — they can simply reuse them, focusing their own limited capacity on learning genuinely new, higher-level combinations. This **direct gradient flow** (each layer has a short path back to the loss through its many connections) and this reuse of already-computed features together allow DenseNet to reach competitive accuracy with markedly **fewer parameters** than architectures relying only on sequential or additive (residual) connectivity.

Q7. Which CNN architecture would you recommend for plant-disease classification?

Evaluated jointly on accuracy potential, computational cost, dataset size, generalization, feature reuse, transfer-learning suitability, and training complexity, **DenseNet (DenseNet-121)** is the most well-rounded recommendation for this scenario, primarily because its parameter efficiency and extensive feature reuse directly address the dominant risk in this problem — overfitting on a small, high-variation dataset — while still offering strong gradient flow for stable training and strong transfer-learning performance from pretrained weights. This is not an unconditional claim that DenseNet is universally best: ResNet-50 remains a very strong, simpler-to-train alternative, particularly if the memory overhead of DenseNet's feature concatenation becomes a practical constraint during training, and the final choice should be confirmed empirically on the actual dataset rather than assumed from architectural analysis alone.

Final Decision Matrix

Note: The scores below are analytical judgments derived from the architectural analysis in this report for the specific plant-disease scenario. They are not measured experimental results.

Criterion	AlexNet	VGGNet	ResNet	DenseNet
Accuracy potential	2/5	3/5	4/5	4/5
Computational efficiency	3/5	2/5	3/5	4/5
Small dataset suitability	2/5	2/5	3/5	4/5
Feature reuse	1/5	1/5	3/5	5/5
Gradient flow	2/5	2/5	4/5	5/5
Generalization	2/5	2/5	4/5	4/5
Transfer learning suitability	3/5	4/5	4/5	4/5
Overall suitability	2/5	2/5	4/5	4/5

FINAL CONCLUSION

This report has analysed the building blocks of CNN-based image classification systems through the lens of a concrete, scenario-driven design problem — plant disease classification from leaf images. Convolution operations play distinct, complementary roles: **standard convolution** provides strong, general-purpose feature extraction but can become computationally expensive at scale; **strided convolution** provides learned down-sampling; **dilated convolution** increases the receptive field without a proportional parameter cost; and **depthwise** and **pointwise** operations, particularly when combined as **depthwise separable convolution**, sharply reduce computation, making them well suited to lightweight, deployable systems.

Regularization techniques address the overfitting risk that arises whenever a high-capacity CNN is trained on a comparatively small, specialized dataset: **data augmentation**, **dropout**, **batch normalization**, **weight regularization**, and **early stopping** each improve generalization through a distinct mechanism, and are most effective when combined.

CNN architecture choice strongly affects both trainability and final performance. **ResNet** solves the degradation problem that limits plain very-deep networks through residual (skip) connections, while **DenseNet** extends this idea to encourage extensive feature reuse through dense connectivity, achieving strong accuracy with markedly fewer parameters.

For the plant-disease classification scenario specifically, the recommended combination balances accuracy, computational efficiency, small-data generalization, and robustness: a **depthwise separable convolutional backbone** (optionally supplemented with dilated convolution for multi-scale context), organised as a **DenseNet-style architecture**, trained with **data augmentation**, **dropout**, and **early stopping**. No single technique in isolation would adequately address the combined challenges of a small dataset and high real-world image variation — it is the *combination*, justified analytically throughout this report, that constitutes the strongest recommended solution.

LEARNING OUTCOMES

Through this comparative analysis task, the student demonstrates the ability to:

- Analyse different convolution operations and their computational and representational trade-offs.
- Compare computational complexity across convolutional and architectural choices.
- Explain the mechanism of hierarchical feature extraction in CNNs.
- Analyse regularization techniques and their distinct contributions to generalization.
- Distinguish underfitting, good generalization, and overfitting from training/validation behaviour.
- Compare AlexNet, VGGNet, ResNet, and DenseNet across multiple architectural dimensions.
- Select appropriate CNN components for a real-world, resource- and data-constrained problem.
- Justify technical design decisions with quantitative and architectural reasoning.
- Analyse trade-offs between accuracy and computational cost across techniques.
- Propose a complete, justified CNN-based solution for plant disease classification.

REFERENCES

- [1] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, "Gradient-based learning applied to document recognition," *Proceedings of the IEEE*, vol. 86, no. 11, pp. 2278–2324, 1998.
- [2] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet classification with deep convolutional neural networks," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 25, 2012.
- [3] K. Simonyan and A. Zisserman, "Very deep convolutional networks for large-scale image recognition," in *International Conference on Learning Representations (ICLR)*, 2015.
- [4] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.
- [5] G. Huang, Z. Liu, L. van der Maaten, and K. Q. Weinberger, "Densely connected convolutional networks," in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.
- [6] N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov, "Dropout: A simple way to prevent neural networks from overfitting," *Journal of Machine Learning Research*, vol. 15, pp. 1929–1958, 2014.
- [7] S. Ioffe and C. Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift," in *International Conference on Machine Learning (ICML)*, 2015.
- [8] F. Chollet, "Xception: Deep learning with depthwise separable convolutions," in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.
- [9] A. G. Howard et al., "MobileNets: Efficient convolutional neural networks for mobile vision applications," *arXiv preprint arXiv:1704.04861*, 2017.
- [10] F. Yu and V. Koltun, "Multi-scale context aggregation by dilated convolutions," in *International Conference on Learning Representations (ICLR)*, 2016.
- [11] D. P. Hughes and M. Salathé, "An open access repository of images on plant health to enable the development of mobile disease diagnostics," *arXiv preprint arXiv:1511.08060*, 2015.
- [12] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. Cambridge, MA, USA: MIT Press, 2016.